

How to Use the CCS Data-Sufficiency Explorer

A step-by-step guide to the interactive dashboard

What this is. An interactive page that answers one question for any of Georgia's continuous count stations: *how many valid days does this station need in this calendar month before its monthly average can be published?*

What you need. A web browser. Nothing to install, no internet connection, no statistics background.

Coverage. All **233** stations, all **12** months — **2,796** station-months — under **six** different outage patterns.

This guide walks through the tool from the first double-click to reading the final number. Everything it describes is done by clicking. Nothing you do in the tool can change or damage the underlying data: the page only reads itself.

Before you start: what this tool is, and what it is not

Georgia DOT asked a research project a single question: when a month at a continuous count station arrives with days missing, how many good days does it actually need before its monthly average can be published? A network-wide rule of 26 days had been applied to every station alike. The project replaced it with an answer worked out separately for each station and each calendar month.

That work is finished. For all 233 stations and all twelve months, the required number of days has already been calculated, checked, and frozen. It is printed in the delivered spreadsheet and in the final report, and it is the figure this page calls the **recommended planning value**.

So what is this page for? Two things.

- **It looks the answer up.** Choose a station and a month and the recommended planning value appears at once. For most users this is the whole job, and the rest of the page is optional.
- **It shows the working.** The tabs underneath let you watch the test that produced that number, see the stations yours was compared against, and check the page's arithmetic against the published figures.

One thing to be clear about from the start. The recommended planning value does not move. You can change the statistical settings on the experiment tab and watch every chart on that tab recompute — that is what it is for — but the recommended value stays exactly as it was certified, at $\alpha = 0.10$, $\epsilon = 0.05$, a 90% success target and 1,000 replications. At other settings the right answer genuinely would be different, but working it out again means re-running the whole selection

experiment, which is far heavier than anything this page does live. If you change a setting, the tool says so on screen, and **Reset to defaults** puts it back.

Contents

Before you start	1
1 The question the tool answers	3
2 Opening it	3
3 Your first run, in five steps	4
4 The screen, region by region	8
5 The controls	9
6 The six outage patterns	12
7 The six answer tiles	13
8 When the month is not complete	14
9 The five tabs	15
9.1 Run an experiment	15
9.2 Planning threshold — where the recommended number comes from	17
9.3 Reference pool — who the comparable stations are	19
9.4 Verification — the page checking itself	20
9.5 How to read this — the method note	21
10 Seeing why the day count matters	21
11 Weekend-heavy: shown, but not recommended	22
12 What the tool does not claim	23
13 If something looks wrong	23
14 Glossary	24

1 The question the tool answers

A continuous count station reports traffic every day. When equipment fails, or a day is rejected in quality control, the month is left partly observed. The practical question follows immediately: *is what survived enough to publish a monthly average?*

Counting the surviving days does not settle it. Consider two Aprils that both kept exactly fourteen days:

- In the first, the fourteen are scattered across the month. Weekdays and weekends are both still represented, so the surviving days still describe the month.
- In the second, the counter died on the 7th and came back on the 24th. The fourteen survivors are two clumps at either end, and everything in the middle of the month is guesswork.

Same count, very different evidence. So the tool asks for two things — how many days survived, *and* how they went missing — and gives back a specific number of days for that combination.

The number is guidance for a *future* month. A month you already hold, and which came in complete, is simply reported: its average is known exactly and no prediction is involved. The recommendation tells you what to aim for next time.

2 Opening it

You will have been sent a folder. Unzip it anywhere you like — the desktop is fine — and inside you will find five things:

File	What it is
GDoT Data-Sufficiency Explorer.html	The tool. Double-click it.
dashboard_how_to_use.pdf	This guide.
practitioner_site_month_kstar.csv	The delivered spreadsheet: the same recommendations as a lookup table, one row per station and month.
practitioner_site_month_kstar_how_to_read.pdf	The guide to that spreadsheet.
README.md	A short plain-text summary of the same thing.

The tool and the spreadsheet carry the same numbers and agree on all **2,796** station-months. Use the spreadsheet to look a value up; use the tool to see the evidence behind it.

Double-click the first one. It opens in your web browser and that is the whole procedure. There is nothing to install, no program to run first, no internet connection needed, and nothing you do leaves your machine.

- Everything the tool needs is inside that single file, including the data for all **233** stations. It is large for that reason, so the first load takes a few seconds. After that it is immediate.

- You can forward that one file to a colleague and it will work for them exactly as it does for you.
- On some computers `.html` files are set to open in a code editor rather than a browser, so you see markup instead of the dashboard. Right-click the file → **Open With** → your browser. You only need to do this once.

If you are holding an earlier copy of the final report. Its dashboard appendix refers to `gdot_dashboard.html` in a `dist` folder. That file has been renamed to **GDoT Data-Sufficiency Explorer.html** and moved to the top of this folder. Same tool, same data, same numbers — the instruction above is the current one.

3 Your first run, in five steps

This section produces a real answer. It uses station 153-0189 in April, which is the case study in the final report, so the numbers here can be checked against it.

Step 1 — Choose the station, the month, and the outage pattern

Everything starts in the control strip along the top.

STATION	CALENDAR MONTH	MISSINGNESS MECHANISM ⓘ
153-0189 — two way, US-129 ▼	April — 30 of 30 valid days ▼	Block-1 ▼


Set **Station** to 153-0189, **Calendar month** to **April**, and **Missingness mechanism** to **Block-1** (one unbroken outage). The three statistical settings — α , ϵ and R — are not in this strip; they sit on the **Run an experiment** tab and affect nothing outside it. Leave them as they ship. Section 5 explains them, and the delivered values are the ones the report used.

Notice that the month is labelled **April — 30 of 30 valid days**. Every month in the list carries its own valid-day count, so you can see what actually arrived before choosing.

Step 2 — Read the six answer tiles

The row of tiles beneath the controls describes the station-month you just chose.

FACILITY ⓘ two way two way · FC 3 · urban · US-129	2024 TARGET STATE ⓘ 30 of 30 valid days — complete every day passed the quality checks	EXACT MADT ⓘ 17,025 vehicles/day, full calendar
---	---	--



Read them left to right: this is a two-way station on US-129; April 2024 arrived **complete**, all 30 of 30 days valid; its true monthly average was **17,025** vehicles/day; its normal scale for April is **15,311**; the recommended planning value is **13** days; and **41 / 43** comparable stations supplied the evidence.

The headline answer is the fifth tile: 13 days. For this station, in April, thirteen valid days are enough to publish a monthly average that is both reliable and tight enough to act on — provided the days that went missing did so in one unbroken stretch.

Step 3 — See which days survive

The first tab, **Run an experiment**, shows the month itself.



The slider sets k , the number of valid days the month is left with; here it is at 13 of 30, which the tool also reports as 43% of the month. Solid blue bars are days that survived, and their height is that day's traffic. Hollow dashed outlines are days that were lost. Because the pattern is Block-1, the lost days form one continuous run — the counter goes down on the 7th and comes back on the 24th.


Three buttons sit beside the slider:

Button	What it does
Run experiment	Re-runs the simulation at the current settings.
New random outage	Draws a different outage of the same shape and the same length, so you can see how much the answer moves from one draw to the next.
Clear hand-edits	Undoes any days you switched by hand and returns to a drawn pattern.

You can also **click any individual day** to switch it on or off. That is how you test an outage you actually experienced, rather than a simulated one.

Step 4 — Read the prediction interval

Directly below the calendar is the answer for the days now surviving.



From those 13 days the tool estimates the monthly average at **16,904** vehicles/day, and puts an interval of **16,256 – 17,553** around it — give or take 648 vehicles/day. Because this April was complete, the true answer is known: **17,025**. The estimate is off by 121 vehicles/day, or -0.71% , and the interval contains the truth.

The row of small badges is the scorecard:

Badge	Meaning
radius $\hat{q}$	How wide the interval is, as a fraction of the station's normal scale. Here 0.0424, i.e. about $\pm 4.2\%$.
meets ϵ	Green when the interval came out tight enough against the target; red when it did not.
precision success	Across all the simulated outages, how often the interval was that tight.
mask-marginal containment	How often the interval contained the true value.
n_{cal}	How many comparable stations the margin was learned from.


Step 5 — Check you are inside the certified range

The grey note at the foot of that panel is the sentence that ties the experiment back to the recommendation:

*“Published planning value for Block-1 at this station-month is $\hat{k}^{\text{LCB}} = 13$. You are testing $k = 13$, which is **at or above** it. This is inside the certified regime.”*

That is the whole workflow in one line. Above the recommended number, the guarantee applies. Below it, it does not — and the tool says so plainly rather than leaving you to work it out. Section 10 shows what going below looks like.

4 The screen, region by region




The whole tool fits on one screen. This view is for orientation — each region is shown full size in the sections that follow.

Region	What it is for
Title bar	The question, three badges recording the scope of the delivery and which years played which role, and the Theme button.
Control strip	Station, month, outage pattern, and three statistical settings. Everything below reacts to these.
Answer tiles	Six facts about the chosen station-month, including the recommended day count.
Tabs	Five views of the same station-month, from a hands-on experiment to the evidence behind the published number.
Working area	Whatever the selected tab shows.

The circled i markers

Every control, tile and panel heading carries a small circled **i**. Hover it — or reach it with the keyboard and pause — and it explains that item in plain language.



Nothing is hidden behind these. They repeat and expand what is already on screen, so the tool can be used without this guide open beside it.

Light or dark

The tool follows whatever your computer is set to, so on a machine in dark mode it opens dark. Every screenshot in this guide is light. Nothing else differs — the same controls in the same places showing the same numbers — so if the colours on your screen do not match the pictures here, that is the reason, and nothing is wrong. The **Theme** button at the top right switches between the two whenever you like. It is not remembered: reopening the file returns to your computer’s own setting.

The tables under the charts

Most panels carry a small grey **Table view** line beneath the picture. Click it and the numbers behind that chart open underneath. Some start open and some start closed, so if a figure quoted in this guide is not on your screen, that line is usually where it is.

5 The controls

The controls sit in two places, and the split is deliberate. What you are looking at is chosen at the top of the page. The statistical settings that affect only the experiment live on the experiment tab itself, so that nothing which moves a chart sits next to the row of values it cannot move.

Across the dashboard

STATION

153-0189 — two way, US-129

▼

CALENDAR MONTH

April — 30 of 30 valid days

▼

MISSINGNESS MECHANISM ⓘ

Block-1

▼

Control	Ships as	What it does
Station	—	Which counting station. All 233 that delivered data in 2024.
Calendar month	—	Which month, labelled with the valid days that actually arrived.
Missingness mechanism	Block-1	The <i>shape</i> of the outage. The most important control, and the one Section 6 is about. It also chooses which curve the Planning threshold tab draws, which is why it stays at the top rather than moving onto the experiment tab with the three settings below.
Reset to defaults	—	Puts everything back. See below.

On the experiment tab

SETTINGS FOR THIS EXPERIMENT ONLY

α (MISCOVERAGE) ⓘ
0.10

ϵ (REL. HALF-WIDTH) ⓘ
0.05

REPLICATIONS R ⓘ
1000

These three sit under **Settings for this experiment only**, on the **Run an experiment** tab. They change nothing outside it, and in particular they do not move the recommended planning value — see the second key box below.

Setting	Ships as	What it does
α (miscoverage)	0.10	How often the interval is allowed to be wrong. At 0.10 it is built to contain the true monthly average at least 90 times out of 100. Lower it for a stronger promise and the interval widens to keep it. <i>This dial sets reliability, not tightness.</i>
ϵ (rel. half-width)	0.05	How tight the answer has to be. At 0.05 the interval may reach no more than 5% above or below the estimate — a 10% window. It is a share of the station’s own normal traffic, which is what lets a quiet county road and a busy interstate be judged on the same footing.
Replications R	1,000	How many times the outage is simulated. Each replication knocks out a different set of days and works the answer out again. More replications give steadier numbers and take longer. <i>The page ships at the same 1,000 the report used, so what you see here matches how the published figure was worked out.</i>

Reliability and tightness are different things, and this is the one idea worth pausing on. An interval can be perfectly reliable and uselessly wide (“the average is somewhere between 5,000 and 40,000” is almost certainly true and completely useless), or narrow and wrong. α governs the first, ϵ the second, and the recommended day count is the smallest number of days that satisfies *both*, often enough.

What the dials move, and what they do not. Changing α , ϵ or the replication count *re-runs the experiment*: the interval, its width, and the success rates below all recompute from the underlying data.

The **recommended planning value** does not follow them. It was certified once, at a fixed set of values, and it keeps them:


$\alpha = 0.10$	$\epsilon = 0.05$	success target = 90%
1,000 replications	seed 20260719	six declared outage patterns

Those are the values behind every number in the delivered spreadsheet and in the report. And it is worth being clear that the recommendation is not *independent* of them — it plainly is not. Loosen α and fewer days would do; tighten ϵ and you would need more.

The reason is cost. Deriving that number means re-running the entire selection experiment: every day count, every outage pattern, a thousand replications each, and then a simultaneous statistical correction across all of those cells at once. That is a far heavier calculation than anything the page does live. So the figure stays as it was certified, at the values in the table above — which is also what makes it match the report and the spreadsheet.

If you do want to explore other settings. Set α and ϵ where you want them, then move the day count and watch the *precision success* figure on the interval panel. That tells you how often a month with that many days would come out tight enough under *your* settings. It is a sound indication, but it is not a certified value: it comes from one fold and carries no correction for the many day counts you tried.

Reset to defaults



Returns the outage pattern, α , ϵ , the replication count, the day count and the random draw to the values the tool ships with, and forgets any days you switched by hand. **Your chosen station and month stay put**, so you can experiment freely and get back to a known state without losing your place.

Two different things this tab does

What the experiment tab shows depends on the month you picked, and the difference matters.

- **If the month arrived complete**, there are no real missing days, so the tool invents some. You are trying out a what-if: this is how the estimate would behave if days were taken away in the pattern you chose. Nothing here is a statement about a month that really happened.

- **If the month really did arrive short**, the tool does not invent anything. It uses the days that actually arrived, and the estimate and range it shows are the operational ones for that month — the tool’s real answer for it.

A banner at the top of the tab says which of the two you are looking at. In neither case does anything on this tab change the recommended planning value above it.

When you have changed something

Move α , ϵ or the replication count away from the delivered values and the tool says so: an amber strip appears under the control row naming what you changed, and the **Planning threshold** tab is marked to warn you that the certified number on it was not computed at your settings. Neither is an error. They are there so a figure read off a changed screen is never mistaken for the delivered one. **Reset to defaults** clears both.

6 The six outage patterns

This is the least familiar control and the one that does the most work. It does not ask *how many* days went missing; it asks *which* ones. The six choices are drawn below. Every panel is the *same* April, and in every panel *exactly fourteen* days survive — only their arrangement changes.

 observed day  missing day

MCAR

Random dropouts

M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

Communication glitches drop isolated days, independently of each other.

Block-1

One outage

M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

A power failure keeps the counter down for one unbroken run of days.

Block-2

Two outages

M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

A modem fault, then a power-supply fault days later: two separate runs.

Weekday-heavy

Working week lost

M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

Weekday road-works: most of the loss lands Mon–Fri, and weekends are hit least.

Weekend-heavy

Weekends lost

M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

Routine weekend maintenance disconnects the sensor every Sat–Sun.

Markov

Bursty failures

M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

A loose connection: once it drops it tends to stay down, and the reverse.

- **MCAR** is the mildest case. The surviving days are spread over the whole month, so they still represent it fairly.
- **Weekend-heavy** is usually the hardest, for the reason set out in Section 11: a month holds only eight or nine weekend dates, and losing them removes a whole block of the weekly rhythm.
- The other four sit in between, and which of them bites hardest depends on the station’s own weekly pattern.

These are deliberate stress tests, not a model of how GDOT equipment actually fails. They exist so a recommendation can be tied to a *stated* assumption instead of an average over unknown failure modes.

Which one should you pick? If you are sizing a station for general planning, use the delivered recommendation, which already takes the worst of five of them. If you know the failure mode you are worried about — a modem that drops out intermittently, a planned closure — pick that pattern and read its number directly.

7 The six answer tiles

Tile	What it tells you
Facility	What is being counted. A two-way station needs both directions reporting; one direction alone is a different quantity and is never used.
2024 target state	What actually arrived. <i>Complete</i> means every day passed quality control, so the average is simply reported. <i>Incomplete</i> means days are missing — the situation this tool is for. <i>No 2024 record</i> means nothing usable arrived for that month at all. Section 8 covers the last two.
Exact MADT	The true monthly average, shown only when the month came in complete, because that is the one case where the answer is known.
Baseline scale B	This station’s normal traffic for this month, fixed from its 2017–2019 history. Errors and interval widths are quoted as a percentage of it.
Recommended planning value	The headline answer: how many valid days to aim for. The small print beneath it says what fraction of the month that is, and what the number would be if weekend-heavy were included — see Section 11.
Cohorts (2022 / 2023)	How many comparable stations supplied the evidence, and from which grouping. Below 9 no usable interval exists at all; 19 or more is preferred for a stable width.

8 When the month is not complete

The walkthrough used April at 153-0189, which arrived whole. Most months did, but not all, and the screen behaves differently for the rest.

First, what “complete” and “incomplete” mean here. They describe the data that arrived, not whether the month can be used. A month is complete when every one of its days produced a valid daily total, and incomplete when at least one did not. A day can fail to produce one for more than a reason: the equipment may have been down, or the day may have arrived and then been set aside in quality control.

An incomplete month is very often still perfectly usable, and saying so is the whole point of the tool. September 2024 at 153-0189 is a fair example: 28 of its 30 days produced a valid total, and the recommended planning value for that station and month is 9 days. The month clears its requirement three times over, and comfortably clears the old 26-day rule as well. The month selector and the status tile therefore show the days themselves — *28 of 30 valid days* — rather than a single word, because the count is what you need and the word was doing more harm than good. Whether that is enough is answered by comparing it with the recommended planning value beside it.

Across the 2,796 station-months in the tool:

2024 target state	Months	What you see
Complete	2,059	The average is known, so it is reported exactly. You can still simulate an outage and watch the interval; because there is a true value, the tool can also tell you whether it was captured.
Incomplete	506	The calendar shows the days that actually arrived rather than a simulated outage, and the estimate shown is the operational one for that month. There is no exact average and no hit-or-miss check, because there is no complete month to check against.
No 2024 record	231	There is no target month to draw at all. The recommended day count is still shown, because it is guidance for a future month and does not depend on this one.


On a month that came in short, the days that never arrived are drawn as flat grey stubs, and clicking one does nothing — there is no reading to switch on. That is not a limitation of the tool. It is the situation a planning threshold exists for: the days are gone, and the question is whether what remains is enough.

The recommendation does not change with the state of the month. It is the same number for a complete April, a half-missing April and an April that never reported, because it says how many valid days that station needs in that calendar month *next* time. What the state changes is how much the tool can show you about the month in hand.


9 The five tabs

9.1 Run an experiment

The hands-on view, and the one Section 3 walked through. Below the calendar and the prediction interval sit two more panels.




How wide is the interval, across replications? The width is not one fixed number — a different set of missing days gives a slightly different margin. This shows the whole spread across the simulated outages, with your ϵ target drawn as a rule. The share of the distribution left of that rule is the precision-success rate quoted on the badge.



Every outage pattern at this k. The same experiment repeated under all six patterns, so you can see which outage shape is hardest for this particular station. No one pattern is always the worst; it depends on the station's weekly rhythm and on how many days are left.

A **Download results as CSV** button at the foot of the tab saves the per-pattern table for the current station-month.

9.2 Planning threshold — where the recommended number comes from




This is the evidence behind the headline number. The curve rises as more days survive: it is the share of simulated outages for which the interval came out tight enough. The dashed rule is the target, $\tau = 0.90$.

The recommendation is not the first point that crosses the target. It is the first point beyond which the curve *stays* above it, all the way up. That distinction is deliberate: a single lucky day count cannot set the number.

Two lines are drawn. The upper one is the raw simulated rate; the lower one is a conservative bound on it, which is what the threshold is actually read from. The gap between them is the price of being sure rather than merely hopeful.

Open **Table view** beneath the curve and the same two lines appear as numbers, one row per day count. *Successes*/*R* is how many of the simulated outages came out tight enough, out of how many were run; *Rate* is that as a fraction, which is the upper line; and *Simultaneous LCB* is the conservative bound, which is the lower one. LCB is short for lower confidence bound.



The second panel puts all six patterns side by side for this station-month. For 153-0189 in April: MCAR needs 8 days, Markov 10, weekday-heavy 11, Block-2 12, Block-1 13 — and weekend-heavy 20. The **Robust (max)** row at the foot of the table reports **13**, the largest of the five ordinary patterns. Section 11 explains why the 20 does not set it.

That table carries two further columns, both of them *descriptions of what happened* rather than promises. *2023 calibration precision* is how often the interval came out tight enough in the year the margin was calibrated, and *2024 target containment* is how often it actually contained the true average in the year that followed. They are worth looking at, and they are *not* the guarantee: containment is measured on the months this station happened to have, not asserted for the next one.

9.3 Reference pool — who the comparable stations are

Resolved reference cohorts

These are the actual stations whose records become conformal scores. They are chosen by the predeclared cluster → road-family → topology hierarchy using only completeness indicators, never traffic outcomes. The target station is always left out.

PREDECLARED SCOPE	2022 PEERS	2023 PEERS	STATUS
cluster selected	41	43	preferred (≥19)
road family	85	90	preferred (≥19)
topology	187	183	preferred (≥19)

The hierarchy takes the **narrowest** scope with at least 19 distinct stations in both roles; failing that, the scope maximising the smaller cohort provided it reaches 9 — the finite-radius floor for 90% split conformal. Cluster id m04:two_way:arterial:n247.

Threshold-selection cohort — 2022

STATION	FACILITY	FC	MADT 2022	BASELINE B
015-0178	two way	3	12,572	10,797
021-0116	two way	4	4,129	4,212
021-0158	two way	3	25,355	29,433
021-0258	two way	3	21,639	25,179
021-0541	two way	4	17,392	18,854
021-0587	two way	4	8,046	8,708
023-0152	two way	4	2,960	3,496
051-0264	two way	3	26,962	22,477
051-0649	two way	4	3,585	3,759
059-0118	two way	4	22,738	24,344
059-0367	two way	4	12,279	12,978
063-1172	two way	4	20,988	20,125
067-0330	two way	3	41,690	45,244

Conformal calibration cohort — 2023

STATION	FACILITY	FC	MADT 2023	BASELINE B
013-0036	two way	4	7,937	8,426
015-0178	two way	3	12,803	10,797
021-0116	two way	4	3,944	4,212
021-0267	two way	4	15,210	15,793
021-0541	two way	4	16,711	18,854
021-0587	two way	4	6,851	8,708
023-0152	two way	4	2,898	3,496
051-0264	two way	3	25,944	22,477
051-0649	two way	4	3,433	3,759
059-0118	two way	4	22,534	24,344
059-0367	two way	4	12,465	12,978
063-1032	two way	4	19,539	19,140
067-0330	two way	3	20,354	24,667

Every margin of error in this tool is learned from how far the method missed at *comparable* stations. This tab names them: the stations used to select the threshold, the stations used to calibrate the interval, and which grouping they were drawn from.

The three scopes at the top

Comparable can be drawn tightly or loosely, so the tool tries three predeclared groupings and reports all of them:

Scope	Who is in it
cluster	The closest grouping: stations of the same kind on the same family of road whose traffic shapes were grouped together in advance.
road family	Wider: the same kind of station on the same family of road, without the finer grouping.
topology	Widest: any station of the same kind.

The two number columns are how many comparable stations that scope actually supplies, in 2022 and in 2023. The rule then picks the **narrowest** scope with at least 19 in both years; if none reaches 19, it takes the one that leaves the larger cohort, provided it reaches 9. Nine is a hard floor — below it no usable interval exists at all — and 19 is preferred because the width steadies. The chosen row is badged *selected*, and the status column says which of the two bars it cleared.

Closing that paragraph is a **cluster id** such as m04:two_way:arterial:n247. It is an internal label for the exact group, readable left to right as the month, the kind of station, the road family,

and an identifier for the grouping itself. You never need it to use the tool; it is there so a result can be traced back to the grouping that produced it.

What the tab guarantees

- **The station is never in its own comparison group.**
- Groups are formed in advance from years older than anything being reported, so a target outcome can never influence its own peer set.
- The route-aligned companion list at the foot of the tab is an engineering diagnostic only. It never decides eligibility, ordering, or stopping, and it is not the comparison group. Its *tier* column says why each station is listed: `lrs_corridor` for the same physical corridor, `posted_route` for the same route number, and `cluster_supplement` for one drawn from the same grouping to fill the list out.

9.4 Verification — the page checking itself

In-browser verification
 This page recomputes the manuscript's arithmetic in JavaScript, so it checks itself against the frozen published results shipped in the bundle. Every check below runs live in your browser right now.

9 of 9 checks pass

Persistent-tail rule reproduces every published K 16776 / 16776 mechanism-months	PASS
Simultaneous Clopper–Pearson bound recomputed in-browser 50,328 sampled cells · max $ \Delta = 5.00\text{e-}7$	PASS
MADT equals the full-calendar mean of the facility-day vector 6,967 complete station-months · max $ \Delta = 5.82\text{e-}11$	PASS
Expanded template has unit mean ($H = 1$) 1,385 peer templates · max $ \Delta = 6.66\text{e-}16$	PASS
Profile-ratio rule on a full mask returns the exact MADT max $ \Delta = 1.25\text{e-}13$	PASS
Conformal radius takes order statistic $t(n+1)(1-\alpha)$ $n = 9\ldots31$	PASS
Finite-radius floor: $n = 8 \rightarrow \infty$, $n = 9 \rightarrow \text{finite}$ $\alpha = 0.10$, matching Proposition 3	PASS
Every mask kernel returns exactly k observed days 6 laws × 3 calendars × 4 values of k	PASS
Every cohort clears the floor, is distinct-station, and excludes the target 2796 station-months checked	PASS

These are exact, deterministic identities — not statistical agreement. If any row fails, the numbers on the other tabs should not be trusted.

This tab re-derives the published arithmetic in your browser and shows it agreeing with the delivered values. These are exact, deterministic identities, not statistical agreement: if any row fails, a number on the page is wrong and should not be used. It is here so the tool can be checked rather than taken on trust — on whatever machine it happens to be opened on.

There are **nine** rows, and the heading above them counts how many passed. *9 of 9 checks pass*, with every row green, is what a healthy copy of the tool shows. Anything else — a red row, or fewer than nine rows listed — is worth reporting before the numbers are used.

9.5 How to read this — the method note

What this tool does, and what it does not claim

The question. A calendar month at a continuous count station is only partly observed. How many valid days are enough before a prediction interval for the full-month MADT is both trustworthy and tight enough to use?

What the dials move, and what stays put. The three statistical settings — α , ϵ and the number of replications — drive the *experiment*. Change one and the prediction interval, its width, and the success rates are all worked out again from the underlying data. The **recommended planning value** does not follow them. It is the certified result of the frozen study, computed at $\alpha = 0.10$, $\epsilon = 0.05$, a **90% success target** and **1,000 replications**, and it is the figure the report and the delivered spreadsheet quote.

This is not because the number is independent of those settings — it plainly is not. Worked example, at station 153-0189 in January: with the delivered settings, 16 observed days give an interval of about $\pm 2.99\%$ of the station's normal traffic, which clears the 5% requirement every time. Tighten the requirement to $\epsilon = 0.02$ and that same 16 days clears it *never* — so the honest recommendation at those settings would be a good deal higher than 16. Re-deriving it properly means re-running the entire selection experiment: every day count, every outage pattern, a thousand replications each, with a simultaneous correction across all of those cells at once. That is far heavier than anything computed live in this page, so the certified value is shipped rather than recalculated.

If you want to reason at other settings, set them and then move the day count, watching the *precision success* figure on the interval panel: it tells you how often a month with that many days would come out tight enough under your settings. Treat it as a sound indication rather than a certified value — it comes from a single fold and carries no correction for the many day counts you tried.

The estimand is the whole facility. A two-way station's daily total needs both travel directions; a ramp or one-way station uses its single registered stream. A partial directional total is a different quantity and is not used.

Two separate targets. α controls *coverage* — the chance the interval contains the true MADT. ϵ controls *precision* — how wide the interval may be, as a fraction of the station's frozen baseline scale. τ is the probability of hitting that width target under the declared mask law, and γ is Monte Carlo familywise error. They are four different numbers and are never interchangeable.

The chronology is what keeps it honest. Profiles, templates, baselines, and peer groups come from 2017–2019. The planning value is certified on 2022 alone. Conformal calibration happens later, on 2023, at the already-frozen k . 2024 is held out. A target outcome can never tune a threshold.

The guarantee. Conditional on the frozen selection and an explicit fixed-site exchangeability assumption for the normalized scores, the interval has finite-sample marginal coverage of at least $1 - \alpha$. That statement is marginal over the declared mask law and is made one site at a time.

What is not claimed. No coverage conditional on one realized outage pattern; no simultaneous guarantee across the network; nothing for an undeclared or adversarial mechanism; and no promise that a narrow interval is therefore a covering one. When you hand-edit the calendar you are asking a *realized-pattern* question, which the theorem deliberately does not answer — the tool reports it as a diagnostic and labels it as such.

A complete month needs no prediction. If the month is strictly complete its MADT is reported exactly. The planning value stays prospective: it tells you what to expect from a *future* delivery of that calendar month.


Bundle generated 2026-08-21T19:48:03Z by dashboard/build_bundle.py from the frozen v2 artifacts behind final_deliverables/final_report.tex. Constants: $\alpha = 0.1$, $\epsilon = 0.05$, $\tau = 0.9$, $\gamma = 0.05$, $R = 1000$. Peer floor 9, preferred 19.

A condensed statement of what the tool does and does not claim, always available beside the data. Section 12 of this guide is the same material in longer form.


10 Seeing why the day count matters

The quickest way to understand the recommendation is to break it. Keep everything else fixed and move the slider.

6 days — below the recommendation



13 days — at the recommendation



	6 days	13 days
Interval half-width	$\pm 1,046$ veh/day	± 648 veh/day
As a share of scale	0.0683	0.0424
Against the $\epsilon = 0.05$ target	misses	meets
Precision success rate	10.2%	86.6%
Verdict	below the certified regime	inside it

Both intervals happen to contain the true value. That is the point worth understanding: the six-day interval is not *wrong*, it is *too wide to be useful*, and it is too wide far too often — it hits the precision target only about one time in ten. The recommendation is the day count at which that stops being a gamble.

11 Weekend-heavy: shown, but not recommended

The recommended planning value is the largest requirement among *five* of the six patterns. Weekend-heavy is deliberately left out of it.

Why

- **The arithmetic is brutal.** A month contains only eight or nine weekend dates. A pattern that removes those first takes away a whole block of the weekly rhythm at once, leaving the estimator nothing to anchor on, and the required day count jumps.
- **It would dominate almost everything.** Weekend-heavy sets the largest of the six values in roughly **three station-months out of five**, sometimes by as much as **twelve days**. At 153–0189 in April it asks for 20 days where the next hardest pattern asks for 13.
- **It is a rare failure mode.** An outage landing squarely and only on weekends is unusual. Sizing every station-month for it would push the recommendation far beyond what ordinary interruptions require, and a target nobody can meet is not useful guidance.

What is *not* done

Nothing is hidden. Weekend-heavy stays fully selectable, with its own curve and its own certified count. Wherever it would raise the number, the recommended-value tile prints that higher figure beside the recommendation — at 153-0189 in April it reads “20 *incl. weekend-heavy*”. If a station genuinely faces weekend risk, such as a recurring weekend closure or planned weekend maintenance, select weekend-heavy and read its number directly.

12 What the tool does not claim

- **It does not produce an interval for a month you already hold.** The recommendation is prospective guidance for a future month of the same name.
- **A recommendation is tied to a stated outage pattern.** It is not a guarantee for an arbitrary pattern of missing days. Hand-editing the calendar is labelled a realized-pattern diagnostic for exactly this reason.
- **The guarantee is per station.** It is invoked one station at a time and is not simultaneous across the network.
- **A complete month is reported exactly.** No prediction replaces it.
- **Precision and coverage are separate.** A narrow interval need not contain the truth, and a covering interval need not be narrow.
- **Comparisons across stations are descriptive.** They do not establish that road class or traffic volume *causes* a requirement.

13 If something looks wrong

Symptom	What it means, and what to do
The recommended value is blank	The station-month was not certified at the simulation resolution used. It is a status, not a zero — treat it as “no recommendation”, not “zero days”.
The page is slow on first open	It carries all its own data. Give it a few seconds; after the first load it is immediate.
You see raw markup instead of the dashboard	The <code>.html</code> file opened in a text editor. Right-click → Open With → your browser.
The numbers moved after you changed a setting	They should: α , ϵ , R and the pattern all change the answer. Press Reset to defaults to return to the delivered figures.
A number disagrees with the delivered planning table	The table and this tool use the same recommendation rule and agree on all 2,796 station-months. If they differ, a setting has been changed from its default.
A row on the Verification tab fails	Stop and report it. Those are exact identities; a failure means a number on the page is wrong.

Symptom	What it means, and what to do
The page opens but the station list is empty	The file did not arrive whole. It is about 25 MB, and a download that stops early leaves a page that looks right and has no data in it. Fetch it again and check the size.

14 Glossary

Term	In plain language
MADT	Monthly Average Daily Traffic — the average of the station’s daily totals across a calendar month. The quantity everything here is about.
k	The number of valid days a month is left with after losses.
Planning value	The recommended k : the fewest valid days at which the interval is reliably tight enough. Written $\hat{k}^{\text{LCB}}$ on screen.
Prediction interval	A range around the estimate, built so that it contains the true monthly average at least $1 - \alpha$ of the time.
α	Miscoverage. How often the interval is allowed to miss. Ships at 0.10.
ϵ	Relative half-width. How wide the interval may be, as a share of the station’s normal scale. Ships at 0.05.
τ	The precision-success target: how often the interval must come out tight enough before a day count is certified. Fixed at 0.90.
Baseline scale B	The station’s normal traffic for that month, fixed from 2017–2019. Widths and errors are quoted against it.
Cohort	The set of comparable stations the margin of error is learned from.
Scope	How widely <i>comparable</i> is drawn: the close <i>cluster</i> , the wider <i>road family</i> , or the widest <i>topology</i> . The narrowest one that supplies enough stations is the one used.
Preferred / finite only	Whether the chosen scope supplied at least 19 comparable stations (preferred, for a steady width) or merely cleared the hard floor of 9 (finite only, usable but read with more caution).
Cluster id	An internal label for the exact grouping a result came from, such as <code>m04:two_way:arterial:n247</code> . Useful for tracing a number, never needed to read one.
Simultaneous bound	The lower of the two lines on the certification curve. It is lowered on purpose so that being wrong at <i>any</i> of the day counts and patterns tested is unlikely, not merely at one of them. On screen this is called the Bonferroni correction.
Mechanism	The assumed shape of the outage — which days went missing, not how many.

Term	In plain language
Conformal	The family of methods used here. Its intervals are built by ranking how far the method missed on past comparable cases, rather than by assuming a model for traffic.